

# Project

## Install dependencies on master and worker nodes

```
TASK [Install k8s] *****
changed: [master] => (item=kubelet)
changed: [node1] => (item=kubelet)
changed: [node2] => (item=kubelet)
changed: [master] => (item=kubeadm)
changed: [node1] => (item=kubeadm)
changed: [node2] => (item=kubeadm)
ok: [master] => (item=kubectl)
ok: [node1] => (item=kubectl)
ok: [node2] => (item=kubectl)

TASK [enable kubelet] *****
changed: [master]
changed: [node1]
changed: [node2]

PLAY RECAP *****
master      : ok=8    changed=7    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
node1      : ok=8    changed=7    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
node2      : ok=8    changed=7    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0

[ec2-user@ip-10-0-1-20 ~]$
```

## Create a single-host Kubernetes cluster

### Initialize the master

```
[ec2-user@ip-10-0-1-20 ~]$ ssh -i "justakey.pem" ec2-user@10.0.2.60
Last login: Wed Aug 24 04:20:24 2022 from ip-10-0-1-20.ap-southeast-1.compute.internal

 _ _ _ _ _
|_| ( _ | _ /
 _| \ | _ | _ |
      Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-10-0-2-60 ~]$ sudo kubeadm init --pod-network-cidr=192.168.0.0/16
[init] Using Kubernetes version: v1.25.0
[preflight] Running pre-flight checks
[WARNING FileExisting-tc]: tc not found in system path
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
```

### Configure kubectl

```
kubeadm join 10.0.2.60:6443 --token 08b7r6.9l0hc9qs6tfq51tc \
--discovery-token-ca-cert-hash sha256:7760720c6887bd331a499c4088099925e0cf5df8ae841
[ec2-user@ip-10-0-2-60 ~]$ mkdir -p $HOME/.kube
[ec2-user@ip-10-0-2-60 ~]$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
[ec2-user@ip-10-0-2-60 ~]$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-10-0-2-60 ~]$
```

## Install Calico in the master node

### Install the Tigera Calico operator and custom resource definitions

```
[ec2-user@ip-10-0-2-60 ~]$ kubectl create -f https://raw.githubusercontent.com/projectcalico/calico/v3.24.0/manifests/tigera-operator.yaml
namespace/tigera-operator created
customresourcedefinition.apiextensions.k8s.io/bgppolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/apiservers.operator.tigera.io created
customresourcedefinition.apiextensions.k8s.io/inagesets.operator.tigera.io created
customresourcedefinition.apiextensions.k8s.io/installations.operator.tigera.io created
customresourcedefinition.apiextensions.k8s.io/tigerastatuses.operator.tigera.io created
serviceaccount/tigera-operator created
clusterrole.rbac.authorization.k8s.io/tigera-operator created
clusterrolebinding.rbac.authorization.k8s.io/tigera-operator created
deployment.apps/tigera-operator created
[ec2-user@ip-10-0-2-60 ~]$
```

## Install Calico by creating the necessary custom resource

```
[ec2-user@ip-10-0-2-60 ~]$ kubectl create -f https://raw.githubusercontent.com/projectcalico/calico/v3.24.0/manifests/custom-resources.yaml
installation.operator.tigera.io/default created
apiserver.operator.tigera.io/default created
[ec2-user@ip-10-0-2-60 ~]$
```

## Confirm the pods are running

Every 2,0s: kubectl get pods -n calico-system

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-795d7b8659-pmdx1	1/1	Running	0	79s
calico-node-w7nxs	1/1	Running	0	79s
calico-typha-86cb4877b6-q4p4r	1/1	Running	0	79s
csi-node-driver-tpjvh	2/2	Running	0	37s

## Join the nodes into the cluster

- Node1

```
[ec2-user@ip-10-0-4-44 ~]$ sudo kubeadm join 10.0.2.60:6443 --token 00b7r6.910hc9qs6tfq51tc --discovery-token-ca-cert-hash sha256:7760720c6887bd331a499c4088099925e0cf5df8ae841fe0734df7d3026e043b
[preflight] Running pre-flight checks
[WARNING FileExisting-tc]: tc not found in system path
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
[ec2-user@ip-10-0-4-44 ~]$
```

- Node2

```

[ec2-user@ip-10-0-1-20 ~]$ ssh -i "justakey.pem" ec2-user@10.0.5.11
Last login: Wed Aug 24 04:20:24 2022 from ip-10-0-1-20.ap-southeast-1.compute.internal

      _ _ _
     _(_)_/
    _(_)_/

Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-10-0-5-11 ~]$ sudo kubeadm join 10.0.2.60:6443 --token 08b7r6.9l0hc9qs6tfq5itc --discovery-token-ca-cert-hash sha256:7760720c6887bd331a499c4088099925e0cf5df8ae841fe0734df7d3026e043b
[preflight] Running pre-flight checks
[WARNING FileExisting-tc]: tc not found in system path
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectll -n kube-system get cn kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
 * Certificate signing request was sent to apserver and a response was received.
 * The Kubelet was informed of the new secure connection details.

Run 'kubectll get nodes' on the control-plane to see this node join the cluster.

[ec2-user@ip-10-0-5-11 ~]$

```

Connect to master node to check if all the nodes are connected

```

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-10-0-2-60 ~]$ kubectll get nodes

```

NAME	STATUS	ROLES	AGE	VERSION
ip-10-0-2-60.ap-southeast-1.compute.internal	Ready	control-plane	18m	v1.25.0
ip-10-0-4-44.ap-southeast-1.compute.internal	Ready	<none>	3m9s	v1.25.0
ip-10-0-5-11.ap-southeast-1.compute.internal	Ready	<none>	88s	v1.25.0

```

[ec2-user@ip-10-0-2-60 ~]$

```

Deploy WordPress on the cluster

```

[ec2-user@ip-10-0-2-60 ~]$ nano kustomization.yaml
[ec2-user@ip-10-0-2-60 ~]$ nano mysql-deployment.yaml
[ec2-user@ip-10-0-2-60 ~]$ nano wordpress-deployment.yaml
[ec2-user@ip-10-0-2-60 ~]$ kubectll apply -k ./
secret/mysql-pass-72mh6dg77t created
service/wordpress created
service/wordpress-mysql created
persistentvolumeclaim/mysql-pv-claim created
persistentvolumeclaim/wp-pv-claim created
deployment.apps/wordpress created
deployment.apps/wordpress-mysql created
[ec2-user@ip-10-0-2-60 ~]$

```

```

[ec2-user@ip-10-0-2-60 ~]$ kubectll get all

```

NAME	READY	STATUS	RESTARTS	AGE
pod/wordpress-76b9d4cc49-bhvnj	1/1	Running	0	32s
pod/wordpress-mysql-6c8bfcc555-lx9tv	1/1	Running	0	32s

```

[ec2-user@ip-10-0-2-60 ~]$ kubectll get services

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	106m
service/wordpress	NodePort	10.108.130.63	<none>	80:30092/TCP	33s
service/wordpress-mysql	ClusterIP	None	<none>	3306/TCP	33s

```

[ec2-user@ip-10-0-2-60 ~]$ kubectll get deployments

```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/wordpress	1/1	1	1	33s
deployment.apps/wordpress-mysql	1/1	1	1	33s

```

[ec2-user@ip-10-0-2-60 ~]$ kubectll get replicaset

```

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/wordpress-76b9d4cc49	1	1	1	32s
replicaset.apps/wordpress-mysql-6c8bfcc555	1	1	1	32s

```

[ec2-user@ip-10-0-2-60 ~]$

```

Create loadbalancer

